The Raggman project stores every PDF Document in Postgres.
Each Document is split into Chunks of about 2000 characters,
and a Chunk never spans two Pages.

Answers are generated by the qwen3:8b model running on Ollama.
Embeddings use qwen3-embedding:0.6b with 1024 dimensions.
A failed Ingestion is retried three times before the Document is marked Failed.